Furever Match

A data product helping kitties everywhere
find their furever home

Alexandra Plassaras

What kind of impact could this really have?

93.8 M

Cats owned in the US (2017)

7.6 M

Cats and Dogs re-enter shelters annually

Case Study: National vs. hyperlocal data

petfinder

+

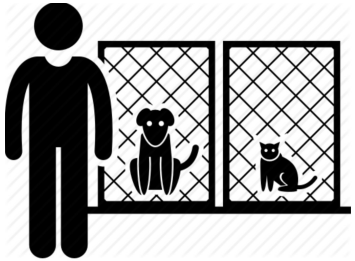

King County

=

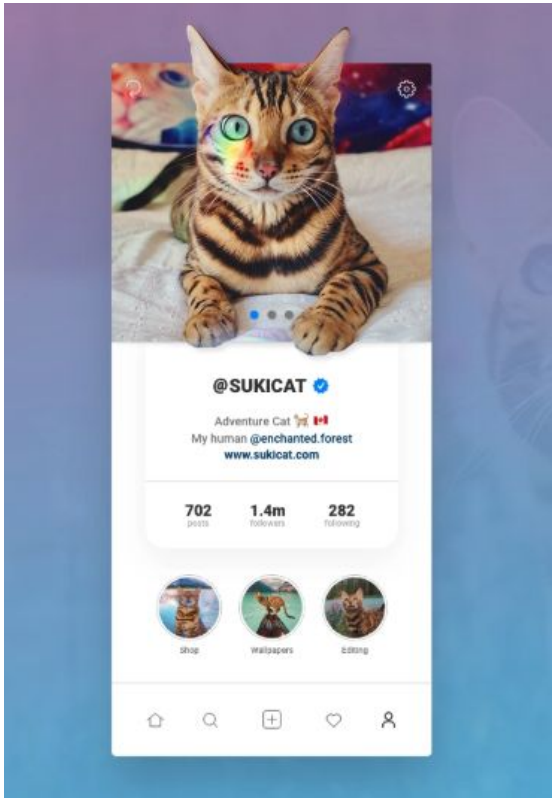

Result: A tool for volunteers, analysts and engineers

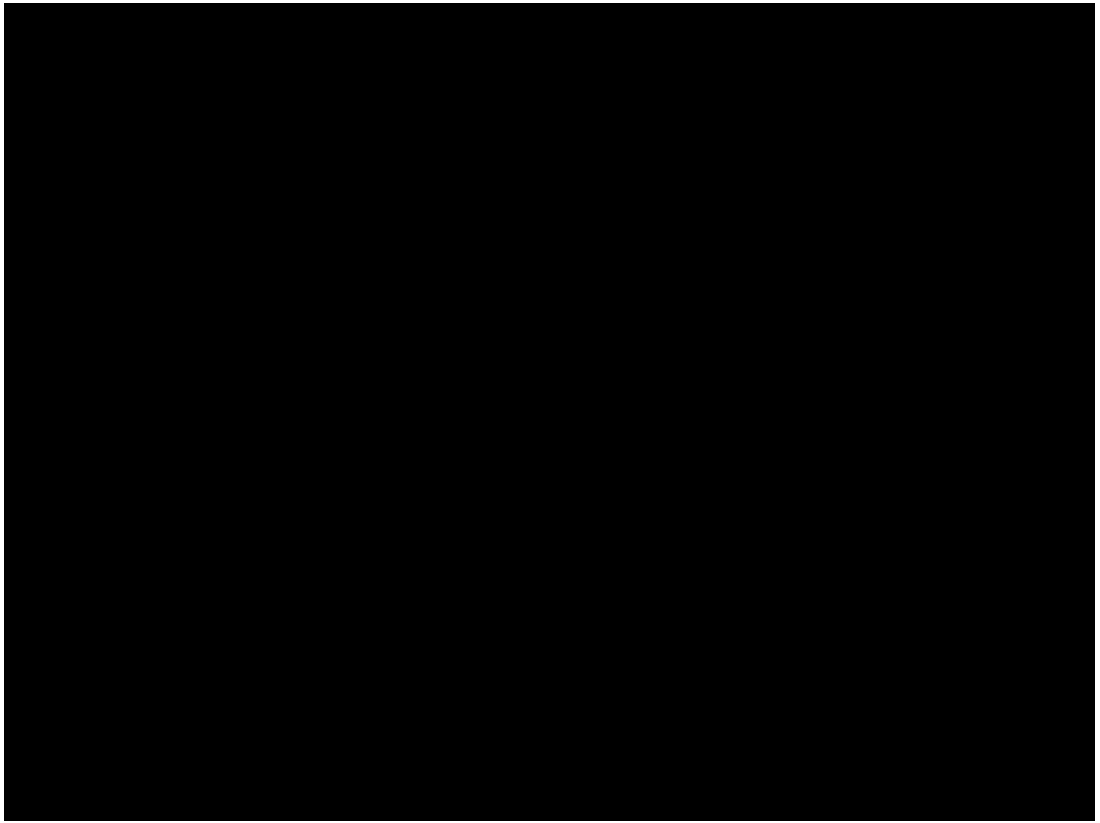

Batch Processing Data Pipeline

Challenge: Acquire data

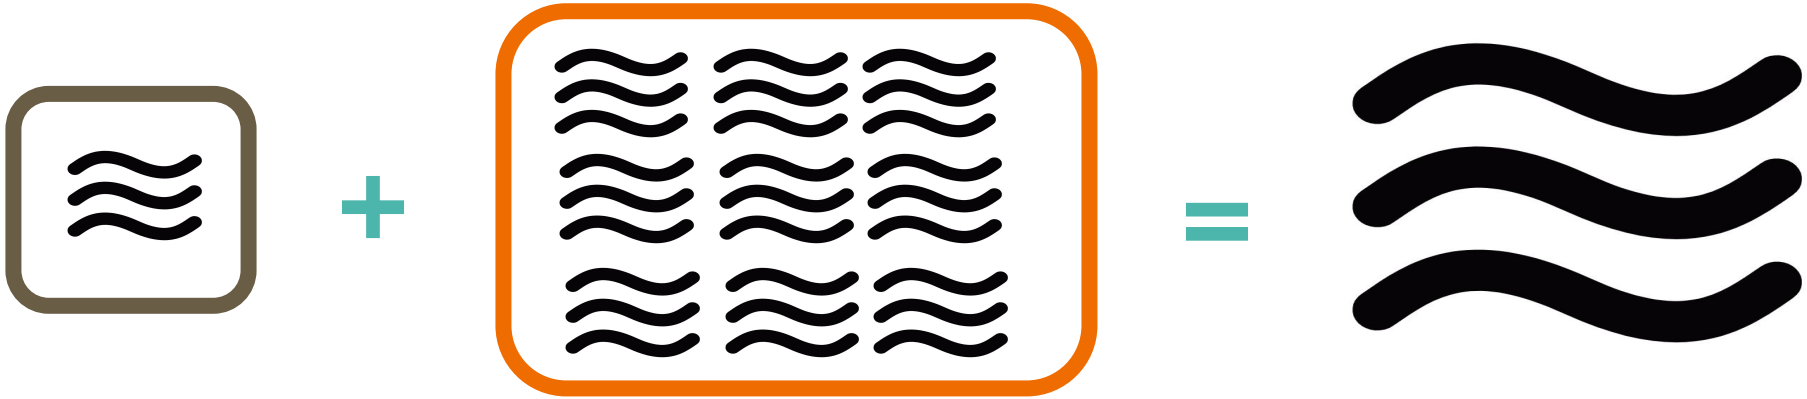

Raw data - Comprised of a subset of states with real cats

Synthetic data - Randomly generated data from all 50 states, top 100 popular cat names

mockaroo

```
pip install pydbgen
```

Challenge: Data Wrangling

local_id	name	age	gender_combo	breed	color	date	date_type	zip	memo	temperament
A606119	Fancy Face	5 YEARS 6 MONTHS	Spayed Female	Domestic Shorthair	Tortie / White	01/25/2020 12:00:...	Received on	98033	Received on: 2020...	BLUE
A606081	Rebel	6 YEARS	Spayed Female	Domestic Shorthair	Tortie	01/24/2020 12:00:...	Received on	98033	Received on: 2020...	RED
A602235	Rose	12 YEARS 1 MONTH	Spayed Female	Domestic Longhair	Black / White	12/09/2019 12:00:...	Received on	98032	Received on: 2019...	BLUE
A576926	Abu	1 YEAR 3 MONTHS	Neutered Male	Domestic Shorthair	Brn Tabby / White	01/26/2020 12:00:...	Received on	98033	Received on: 2020...	GREEN
A606544	Willy	10 YEARS	Neutered Male	Domestic Longhair	Brn Tabby	01/30/2020 12:00:...	Received on	98033	Received on: 2020...	GREEN
A577312	Pepper	7 YEARS 1 MONTH	Spayed Female	Domestic Shorthair	Black	12/21/2018 12:00:...	Received on	98032	Received on: 2018...	BLUE
A598970	Ms Jazz	2 YEARS 3 MONTHS	Spayed Female	Domestic Shorthair	Brn Tabby / White	10/18/2019 12:00:...	Received on	98032	Received on: 2019...	GREEN
A574929	Georgie	9 YEARS 2 MONTHS	Spayed Female	Domestic Shorthair	Brn Tabby / White	01/28/2020 12:00:...	Received on	null	Received on: 2020...	GREEN
A606083	Miley	3 YEARS 7 MONTHS	Spayed Female	Domestic Shorthair	Black	01/24/2020 12:00:...	Received on	98033	Received on: 2020...	RED
A600637	Kuni	7 MONTHS	Neutered Male	Domestic Shorthair	Brn Tabby	11/12/2019 12:00:...	Received on	98032	Received on: 2019...	BLUE

*

- Age is a string with both years and months
- Gender is gender plus spayed/neutered
- Temperament is on a scale compared to good with children, cats and dogs

Alexandra Plassaras

M.S. in Quantitative Methodologies for Social Science

COLUMBIA UNIVERSITY

Udemy

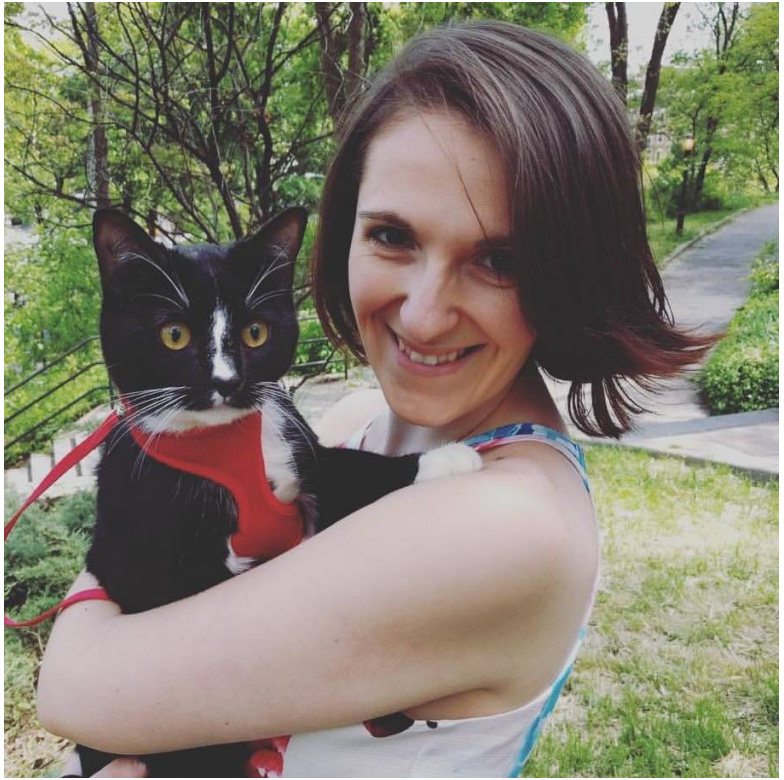

aplassaras@gmail.com

What's next for furever match?

- Build out backend
  - Airflow to automate data ingestion and prioritize data
  - Store historical snapshots
- Bring in the Data Analysts / Scientists
  - NLP for textual data
  - Image processing against adoptable and lost pets
  - Recommendation algorithms for owner + cat

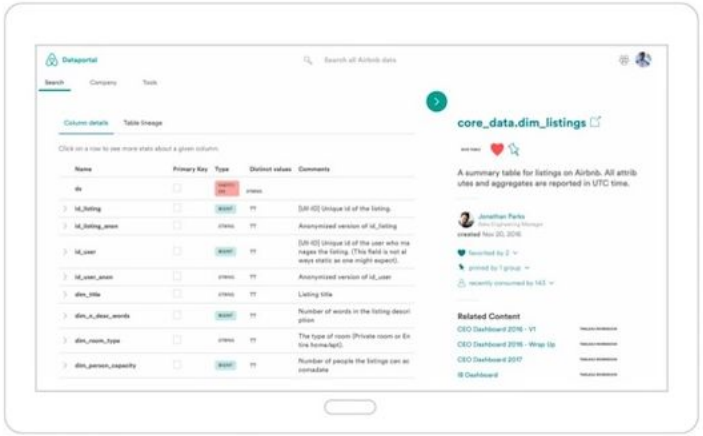

Column details & value distributions
Table lineage
Enrich metadata on the fly

Why I used S3

- Easily integrate with current pipeline and with other AWS services
- Durability of 99.999999999% of objects per year
- Data on S3 persists compared to HDFS which doesn't persist once an instance is stopped
- With S3 you only pay for the storage that you actually need plus S3 compresses files

Why I used Spark?

- Faster than Hadoop
  - Spark tries to do as many calculations as possible in memory, which avoids moving data back and forth across a cluster
- Apache Spark is known for batch-processing big data
- It is open source analytics platform for large-scale processing of huge datasets, large online community
- It has resilient distributed datasets (RDDs), and the in-memory data structure allows Spark to perform functional programming.
- It uses a DAG scheduler along with physical execution engine and the query optimizer.
- It contains a stack of libraries including Spark SQL

Why I used RDS PostgreSQL (compared to Redshift)

- Scaling:
  - Takes only a few minutes for RDS (reconfiguring virtual instances)
- Storage Capacity:
  - I don't need a storage capacity of up to 2 PB
- Data Replication:
  - Don't need to copy complete data to S3 and then copy (Redshift).
  - Depends on the underlying database I use in RDS
- Pricing:
  - Limited budget, RDS already included in current AWS services ($0.017 vs $0.25)
- Performance:
  - RDS has better performance for queries that don't tests its limits (millions of rows)
    - Given my budget for AWS and $0 budget for APIs Isn't an issue for now
- What I could do in the future:
  - Once my data is sufficiently large I could add Redshift to my robust analytical pipeline

Challenge: Acquire data

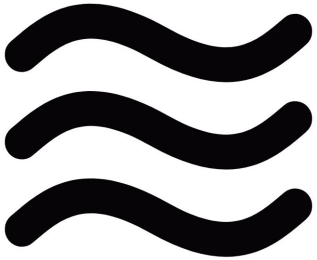

Raw data

Duplicate data

Synthetic data

JSON

Animal info

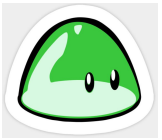

Animal Description

Organization Info

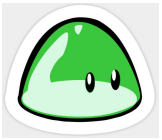

Temperament

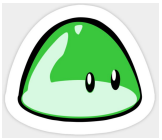

Medical Info

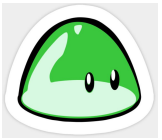

Adoption Status